



Jack Linke

@JackDLinke

jacklinke.com

Managing Director **Watervize** - SaaS for Irrigation



Husband



Public Speaker



ND



Veteran



he/him

# Home on the range with Django

Getting comfortable with ranges and range fields

<https://github.com/jacklinke/home-on-the-range-with-django/>

# Disclaimers

- Topic might be challenging for beginners
  - (But exposure to these concepts early may help avoid pitfalls later)
- Postgres-specific

# Goals

- Learn why ranges can be more useful than distinct start & end values
- Become familiar with range / interval terminology
- See a number of approaches to using & querying with ranges
- Have resources for further reading & learning

# Real Life

- Daily Temperature highs/lows
- Budget range for home renovation
- Range of frequencies for each musical instrument
- Range of gas prices in each region
- Range of calories in fast food meals
- Range of goals/touchdowns for each team & season
- Daily business hours
- Salary range for different jobs
- Suitable soil pH values for different plants
- Start & end date/datetime of Events
- Typical min/max driving duration between locations

# History

- First available in Postgres 9.2 (2011)
- First available in Django in 2015

# Typical Approach

- Separate **start** and **stop** model fields
- Querying with start and stop values
- Quickly gets complicated

## Model Constraint Comparison - separate fields

- 2 separate fields (DB doesn't know they are related)
- Have to build logic to ensure you don't end up with...
  - salary\_range = \$200,000 to \$140,000 😨
  - event\_dates = 2022/10/21 to 2022/10/16 🤔
- Constraint to prevent overlaps
  - Have to write an expression to trick the DB into treating 2 fields (and boundary) as a range type

# Model Constraint Comparison - separate fields

```
1 class TsTzRange(Func):
2     function = "TSTZRANGE"
3     output_field = DateTimeRangeField()
4
5
6 class LockerReservation(models.Model):
7     period_start = models.DateTimeField()
8     period_end = models.DateTimeField()
9     ...
10
11 class Meta:
12     constraints = [
13         # No Locker should have overlapping reservations
14         ExclusionConstraint(
15             name="exclude_overlap_locker_res",
16             expressions=(
17                 (
18                     TsTzRange("period_start", "period_end", RangeBoundary()),
19                     RangeOperators.OVERLAPS,
20                 ),
21                 ("locker", RangeOperators.EQUAL),
22             ),
23             condition=Q(cancelled=False),
24         ),
25         # Make sure the start value is less than or equal to the end value
26         models.CheckConstraint(
27             name="prevent_start_gt_end",
28             check=~Q(period_start_lte=F("period_end")),
29         ),
30         ...,
31     ]
```

## Model Constraint Comparison - range field

- 1 field stores the lower, upper, and boundary information
  - DB knows these are all related
- By default, Postgres won't let you use larger value in lower position
  - `salary_range = $140,000 to $200,000` 😂
  - `event_dates = 2022/10/16 to 2022/10/21` 😎

# Model Constraint Comparison - range field

```
1 class LockerReservation(auto_prefetch.Model):
2     period = DateTimeRangeField()
3     ...
4
5     class Meta:
6         constraints = [
7             # No Locker should have overlapping reservations
8             ExclusionConstraint(
9                 name="excl_overlap_locker_res",
10                expressions=[
11                    ("period", RangeOperators.OVERLAPS),
12                    ("locker", RangeOperators.EQUAL),
13                ],
14                # Ignore overlaps where the reservation is cancelled
15                condition=Q(cancelled=None),
16            ),
17            ...,
18        ]
```

## Query Comparison - separate fields

- Queries have to include both fields
- Often unintuitive in which order lookups are used

## Query Comparison - separate fields

```
1 now = timezone.now()
2
3 next_12_hours_start = now
4 next_12_hours_end = now + timezone.timedelta(hours=12)
5
6 upcoming_events = Event.objects.filter(
7     event_start__lt=next_12_hours_end,
8     event_end__gte=next_12_hours_start
9 )
```

## Query Comparison - separate fields

```
1 now = timezone.now()
2
3 next_12_hours_start = now
4 next_12_hours_end = now + timezone.timedelta(hours=12)
5
6 upcoming_events = Event.objects.filter(
7     event_start__lt=next_12_hours_end,
8     event_end__gte=next_12_hours_start
9 )
```

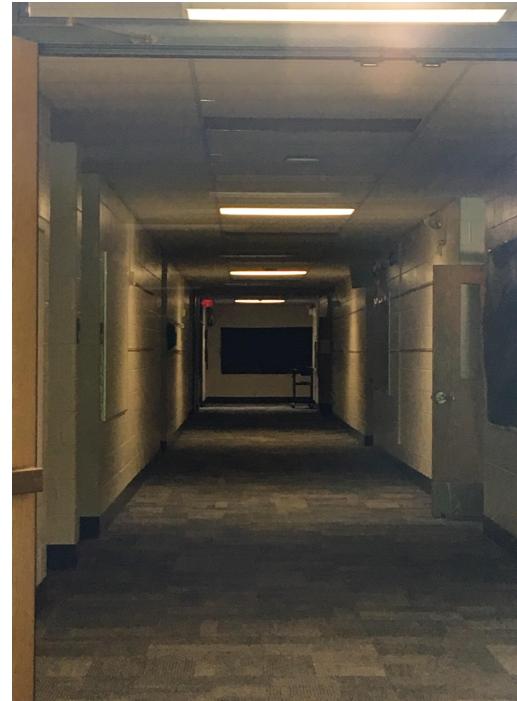
## Query Comparison - range field

- Use a single, intuitive lookup

# Query Comparison - range field

```
1 now = timezone.now( )
2
3 next_12_hours = DateTimeTZRange(
4     now,
5     now + timezone.timedelta(hours=12)
6 )
7
8 upcoming_events = Event.objects.filter(
9     event_period__overlap=next_12_hours
10 )
```

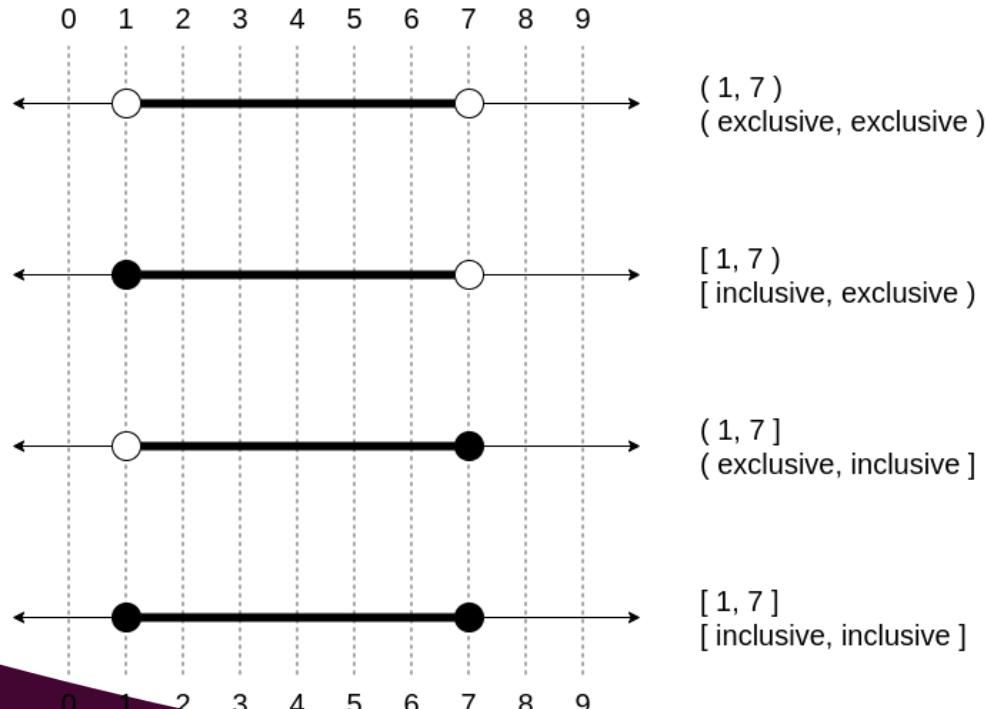
# Going back to Elementary School (but no tests - I promise)



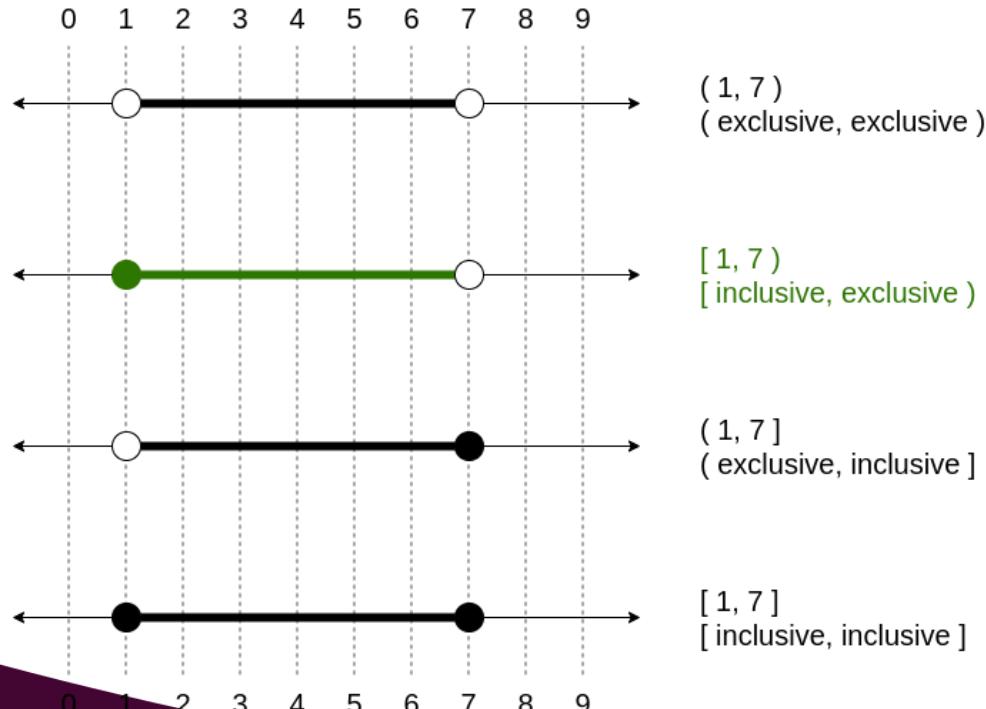
# Terminology

- Ranges = Intervals
- Inclusive = Closed
- Exclusive = Open

# Inclusive vs Exclusive



# Inclusive vs Exclusive



# Inclusive vs Exclusive

Another way to look at this

$$[1, 3] = 1 \leq x \leq 3$$

$$[1, 3) = 1 \leq x < 3$$

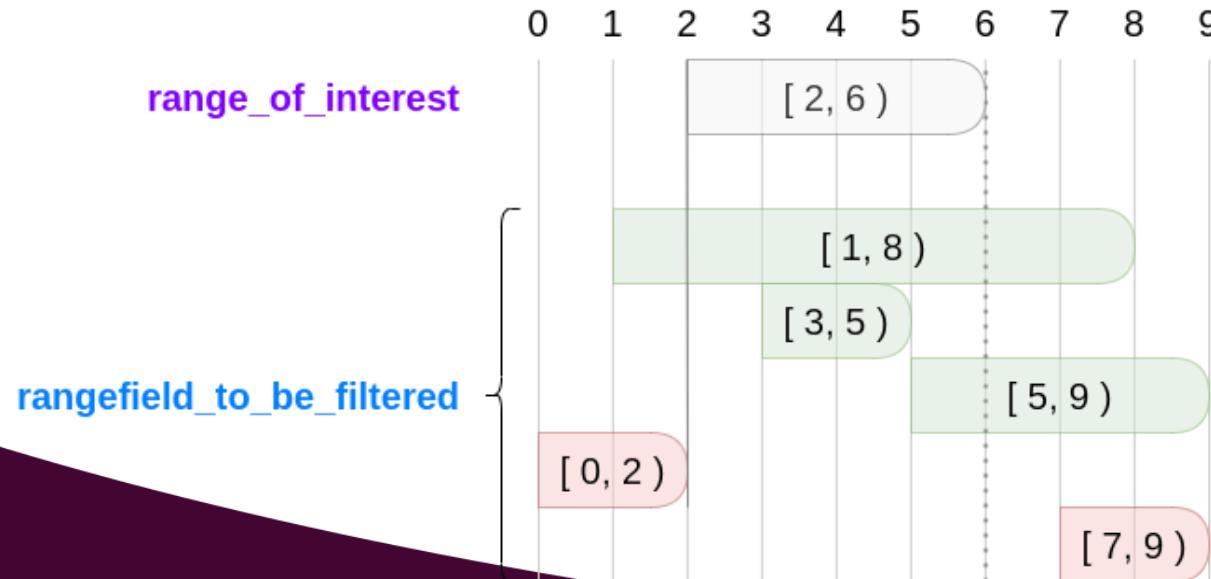
$$(1, 3] = 1 < x \leq 3$$

$$(1, 3) = 1 < x < 3$$

# Range Overlaps Range

`rangefield_to_be_filtered_overlap=range_of_interest`

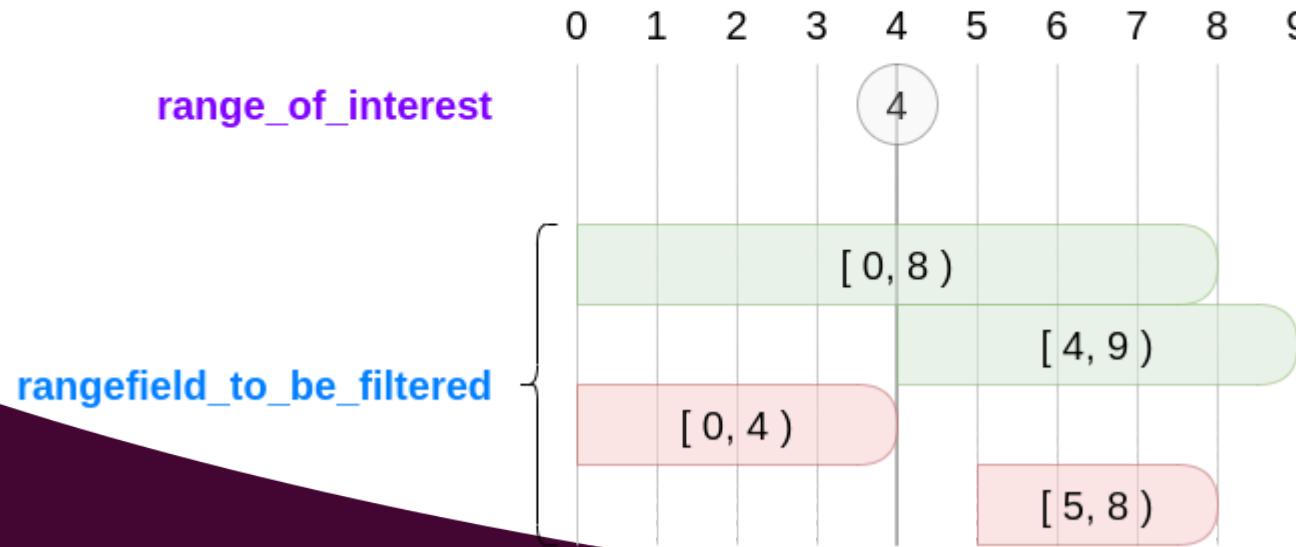
*rangefield\_to\_be\_filtered overlaps range\_of\_interest (by any amount)*



# Range Contains Element

`rangefield_to_be_filtered__contains=element_of_interest`

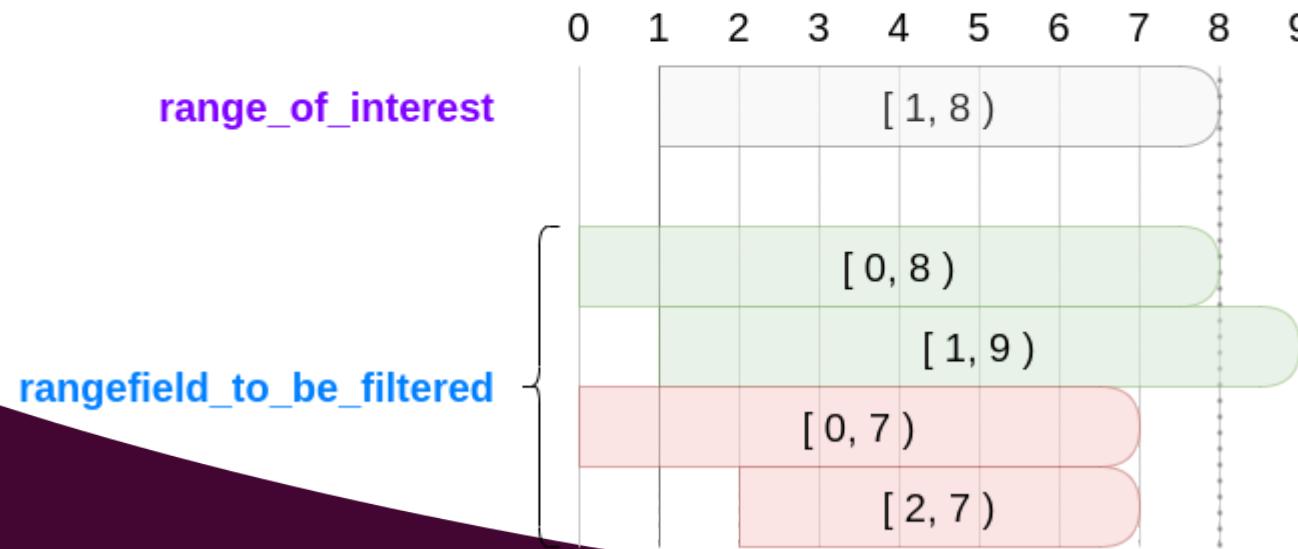
*rangefield\_to\_be\_filtered contains element\_of\_interest*



# Range Contains Range

`rangefield_to_be_filtered__contains=range_of_interest`

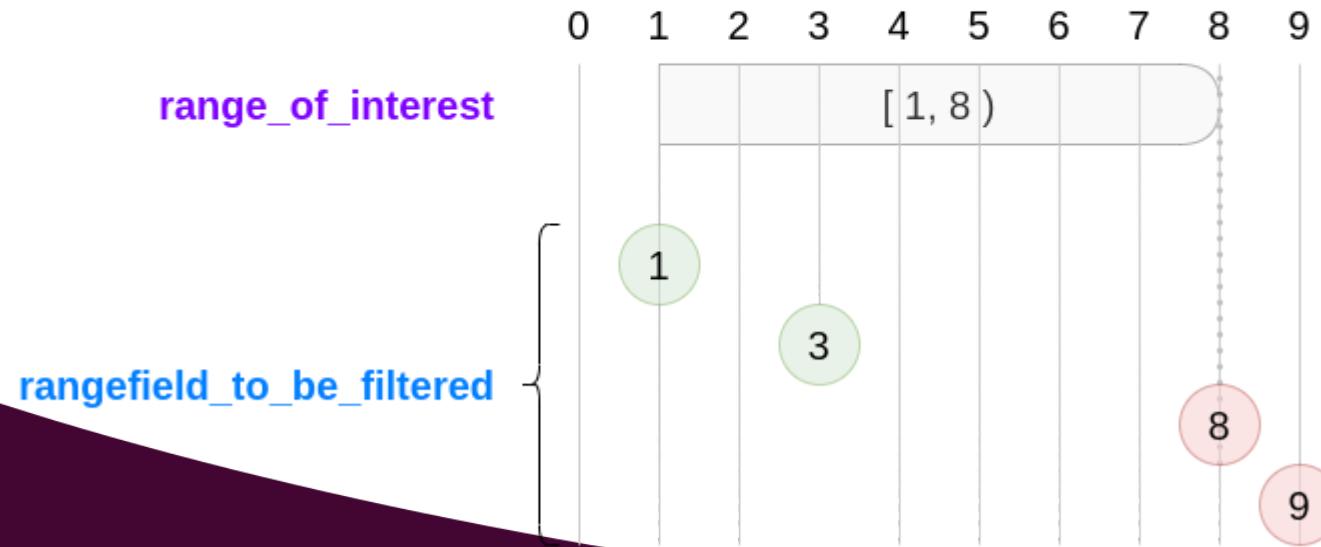
*rangefield\_to\_be\_filtered contains range\_of\_interest*



# *Element Contained By Range*

`field_to_be_filtered__contained_by=range_of_interest`

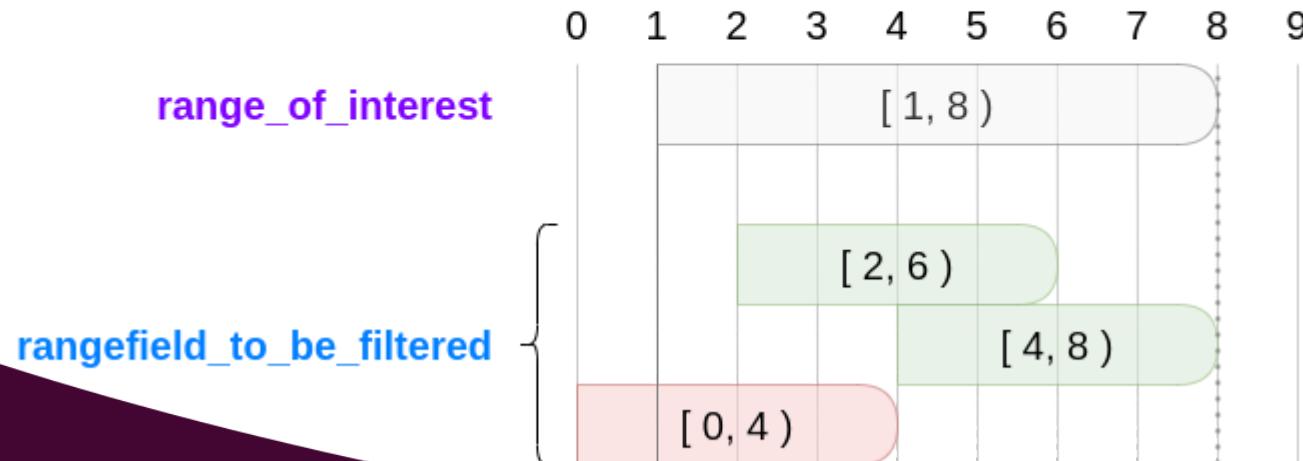
*field\_to\_be\_filtered* is contained by *range\_of\_interest*



# Range Contained By Range

`rangefield_to_be_filtered__contained_by=range_of_interest`

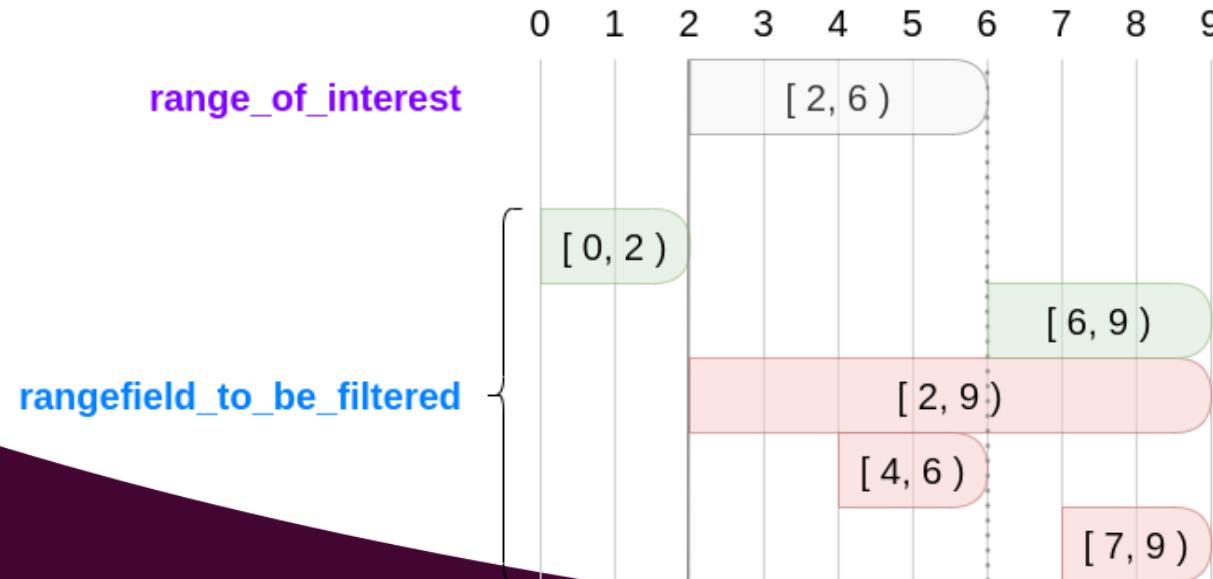
*rangefield\_to\_be\_filtered* is contained by *range\_of\_interest*



# Range Adjacent To Range

`rangefield_to_be_filtered__adjacent_to=range_of_interest`

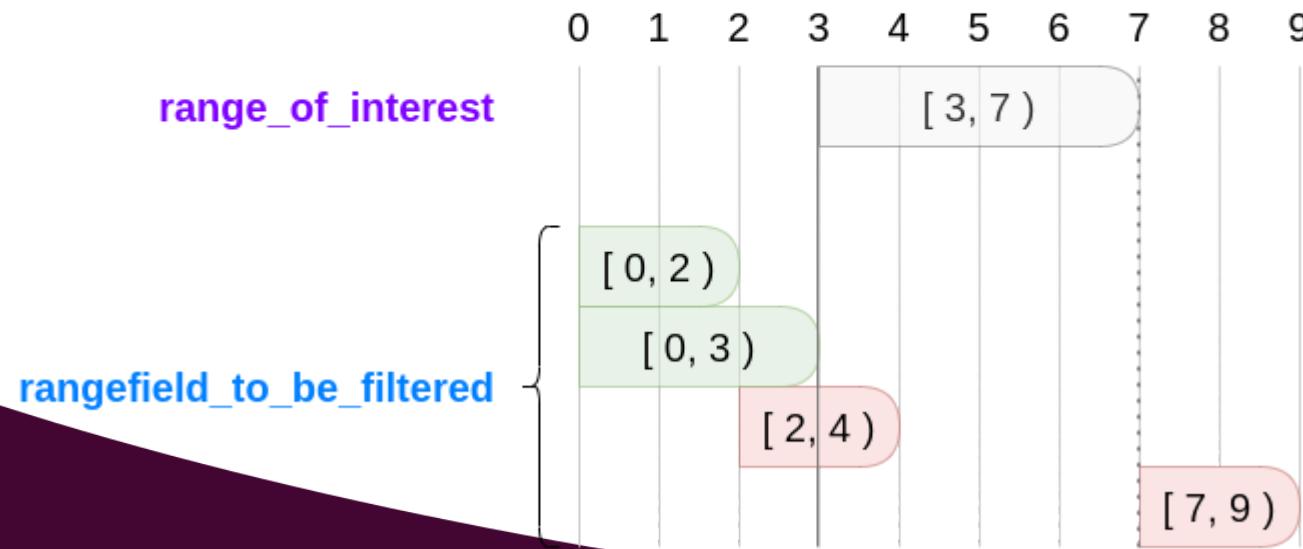
*rangefield\_to\_be\_filtered* is immediately right or left of *range\_of\_interest*



# *Range Fully Less Than Range*

`rangefield_to_be_filtered__fully_lt=range_of_interest`

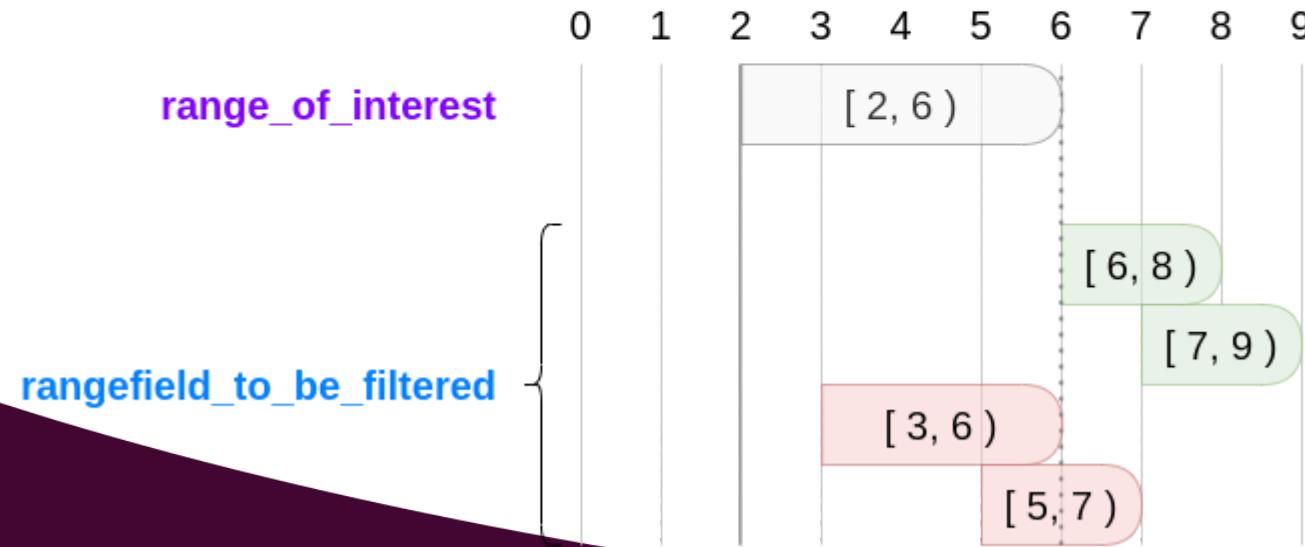
*rangefield\_to\_be\_filtered* is completely to the left of *range\_of\_interest*



# *Range Fully Greater Than Range*

`rangefield_to_be_filtered_fully_gt=range_of_interest`

*rangefield\_to\_be\_filtered* is completely to the right of *range\_of\_interest*



## `psycopg2.extras`

Provides tools for 'translation'

- `NumericRange`

- Postgres: `int4range`, `int8range`, `numrange`
  - Django: `IntegerField`, `BigIntegerField`, `DecimalRangeField`

- `DateRange`

- Postgres: `daterange`
  - Django: `DateRangeField`

## `psycopg2.extras` (cont'd)

Provides tools for 'translation'

- `DateTimeRange`
  - Postgres: `tsrange`
  - Django: `DateTimeRangeField` (naive dt 💩)
- `DateTimeTZRange`
  - Postgres: `tstzrange`
  - Django: `DateTimeRangeField`

# Defining Range Fields in Django

- Using `psycopg2.extras` classes
  - Bounds default to "[]"
  - Can be changed for non-discrete range fields with `default_bounds`
  - (`DateTimeRangeField` and `DecimalRangeField`)
- Fields also accept tuples as input if no bounds information is necessary

```
1 my_field = IntegerRangeField()
2 my_field = IntegerRangeField(default=(1, 3))
3 my_field = IntegerRangeField(NumericRange(None, 7, "( )"))
4 my_field = IntegerRangeField(default_bounds="[]")
```

## Lookups based on range bounds

- **startswith**: filters on lower bound
- **endswith**: filters on upper bound
- **isempty**: filters on empty instances
- **lower\_inf**: unbounded (infinite) or bounded lower bound
  - e.g.: date.min
- **upper\_inf**: unbounded (infinite) or bounded upper bound
  - e.g.: date.max
- **lower\_inc**: filters on inclusive or exclusive lower bounds
- **upper\_inc**: filters on inclusive or exclusive upper bounds

# Lookups based on range bounds

```
1 from psycopg2.extras import NumericRange
2
3 # From -infinity to infinity
4 my_first_range = NumericRange(lower=None, upper=None, bounds='[ )')
5
6 # From 1 (inclusive) to 5 (exclusive)
7 my_second_range = (1, 5)
8
9 # From 7 (inclusive) to infinity
10 my_third_range = NumericRange(lower=7, upper=None)
11
12
13 # In the model
14 class MyModel(models.Model):
15     my_integer_range = IntegerRangeField()
16
17 # In a query
18 first = MyModel.objects.create(my_integer_range=my_first_object)
19 second = MyModel.objects.create(my_integer_range=my_second_range)
20 third = MyModel.objects.create(my_integer_range=my_third_range)
```

# Lookups based on range bounds

```
1 my_first_range = NumericRange(lower=None, upper=None, bounds='[ )')
2 my_second_range = (1, 5)
3 my_third_range = NumericRange(lower=7, upper=None)
4
5 first = MyModel.objects.create(my_integer_range=my_first_object)
6 second = MyModel.objects.create(my_integer_range=my_second_range)
7 third = MyModel.objects.create(my_integer_range=my_third_range)
8
9
10 MyModel.objects.filter(my_integer_range__startswith__gt=3)
11 # <QuerySet [<MyModel: my_third_range]>
12
13 MyModel.objects.filter(my_integer_range__endswith=5)
14 # <QuerySet [<MyModel: my_second_range]>
```

## Creating your own range types

*"every data type that supports a btree operator class (that is, a way of ordering any two given values) can be used to create a range type."*

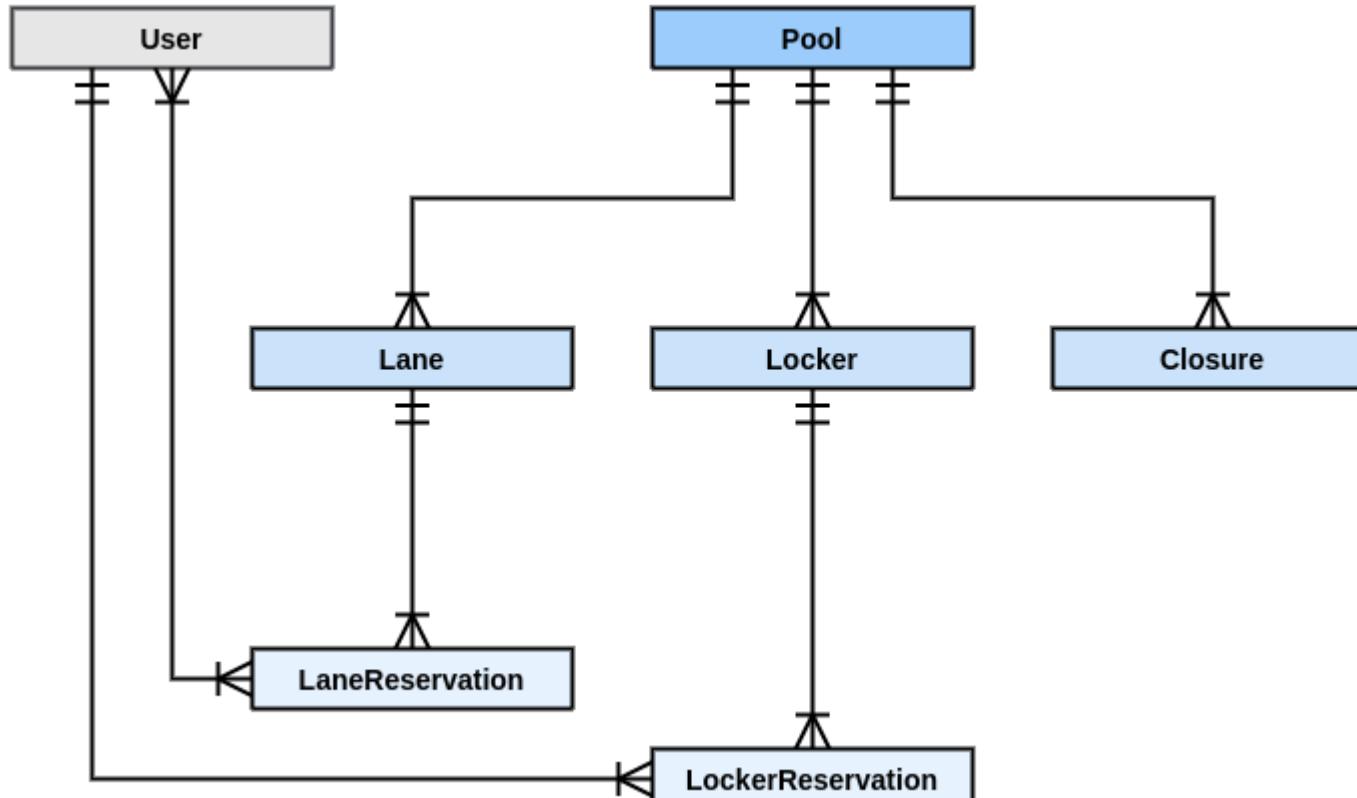
- PostgreSQL 14 Administration Cookbook

*Example: GenericIPAddressField (postgres' inet) can be extended to create range of IP Addresses*

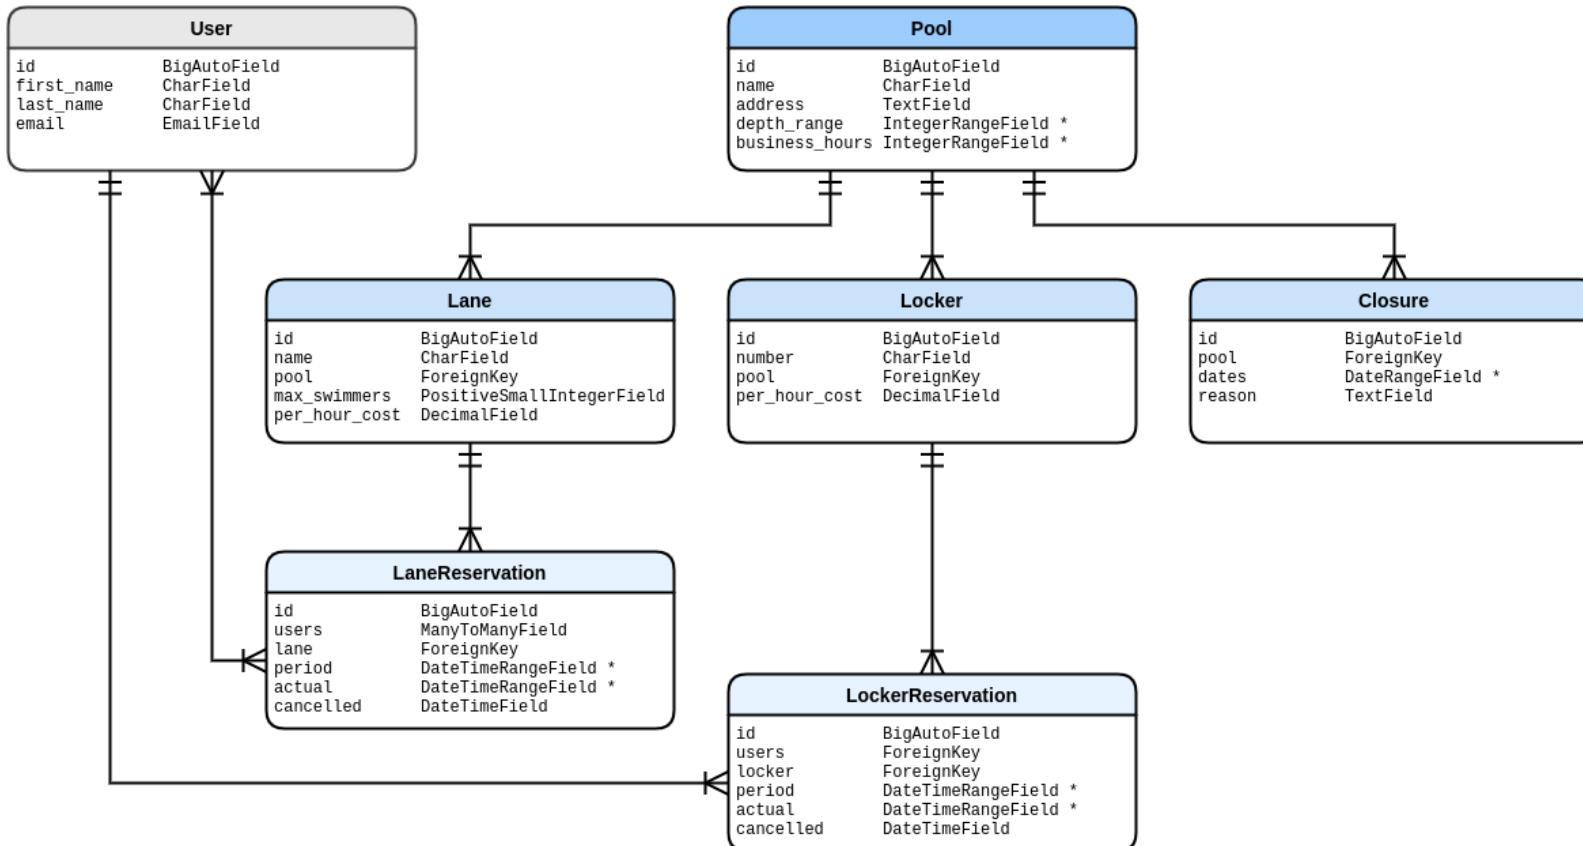
# Indexing

Stick with GiST or B-Tree

# Example Project



# Example Project



# Pitfalls

## DateTimeRangeField in Admin

Reservation Period:

2022-10-15 11:00:00

2022-10-15 13:00:00

Actual Usage Period:

## Pitfalls

Trying to use lower and upper as callables

- They aren't.
- You can use Lower and Upper database functions in queries
- You can use startswith and endswith lookups

## Pitfalls

Trying to use `F()` with the `psycopg2.extra` classes

They aren't Django classes! May have to use `Cast` to convert to a  
Django range field

## Resources

- [psycopg2.extras](#)
- [This talk & Example Project](#)
- [django-range-merge](#) - use `range_merge` aggregate with Django
- [django-generate-series](#) - create sequences with Django's ORM

## Final thought

A man works hard to name an interval equal to 24 hours...  
...so he calls it a day.

